

Exercise2 Pandas

April 14, 2018

1 Hochschule Bonn-Rhein-Sieg

2 Learning and Adaptivity, SS18

3 Assignment 01 (15-April-2018)

3.1 Sathiya Ramesh, Pradheep Krishna Muthukrishnan Padmanabhan, Naresh Kumar Gurulingan

4 Pandas

Pandas is an open source library providing high-performance, easy-to-use data structures and data analysis tools for the Python.

Library documentation: <http://pandas.pydata.org/>

4.0.1 General

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

4.1 Task 1

Create dataframe (that we will be importing)

```
In [2]: data = {'first_name': ['Jason', 'Molly', 'Tina', 'Jake', 'Amy'],
               'last_name': ['Miller', 'Jacobson', ".", 'Milner', 'Cooze'],
               'age': [42, 52, 36, 24, 73],
               'preTestScore': [4, 24, 31, ".", "."],
               'postTestScore': ["25,000", "94,000", 57, 62, 70]}
```

```
In [3]: df = pd.DataFrame(data)
df.head()
```

```
Out[3]:
```

	age	first_name	last_name	postTestScore	preTestScore
0	42	Jason	Miller	25,000	4
1	52	Molly	Jacobson	94,000	24

2	36	Tina	.	57	31
3	24	Jake	Milner	62	.
4	73	Amy	Cooze	70	.

4.2 Task 2

- Save dataframe as csv
- Load a csv
- Load a csv with no headers
- Load a csv while specifying column names
- Load a csv while skipping the top 3 rows

```
In [4]: SAVE_PATH = './task1.csv'
df.to_csv(SAVE_PATH, index=False)
df_task2 = pd.read_csv(SAVE_PATH)
df_task2.head()
```

```
Out[4]:
```

	age	first_name	last_name	postTestScore	preTestScore
0	42	Jason	Miller	25,000	4
1	52	Molly	Jacobson	94,000	24
2	36	Tina	.	57	31
3	24	Jake	Milner	62	.
4	73	Amy	Cooze	70	.

```
In [5]: df_no_headers = pd.read_csv(SAVE_PATH, header=None)
df_no_headers.head()
```

```
Out[5]:
```

	0	1	2	3	4
0	age	first_name	last_name	postTestScore	preTestScore
1	42	Jason	Miller	25,000	4
2	52	Molly	Jacobson	94,000	24
3	36	Tina	.	57	31
4	24	Jake	Milner	62	.

```
In [6]: df_specific_columns = pd.read_csv(SAVE_PATH, usecols= ['first_name', 'last_name'])
df_specific_columns.head()
```

```
Out[6]:
```

	first_name	last_name
0	Jason	Miller
1	Molly	Jacobson
2	Tina	.
3	Jake	Milner
4	Amy	Cooze

```
In [7]: df_skip_3rows = pd.read_csv(SAVE_PATH, skiprows= 3, header= 0)
df_skip_3rows.head()
```

```
Out[7]:
```

	36	Tina	.	57	31
0	24	Jake	Milner	62	.
1	73	Amy	Cooze	70	.

5 It is interesting to know and play around

```
In [8]: # create a series
s = pd.Series([1,3,5,np.nan,6,8])

In [9]: # create a data frame
dates = pd.date_range('20130101',periods=6)
df = pd.DataFrame(np.random.randn(6,4),index=dates,columns=list('ABCD'))

In [10]: # another way to create a data frame
df2 = pd.DataFrame(
    { 'A' : 1.,
      'B' : pd.Timestamp('20130102'),
      'C' : pd.Series(1,index=list(range(4)),dtype='float32'),
      'D' : np.array([3] * 4,dtype='int32'),
      'E' : 'foo' })

df2

Out[10]:
```

	A	B	C	D	E
0	1.0	2013-01-02	1.0	3	foo
1	1.0	2013-01-02	1.0	3	foo
2	1.0	2013-01-02	1.0	3	foo
3	1.0	2013-01-02	1.0	3	foo

```
In [11]: df2.dtypes

Out[11]: A          float64
B    datetime64[ns]
C          float32
D           int32
E           object
dtype: object

In [12]: df.head()

Out[12]:
```

	A	B	C	D
2013-01-01	0.880059	0.477660	0.651964	0.525817
2013-01-02	-0.070696	0.811440	1.042583	1.429121
2013-01-03	-0.417656	1.171581	0.886486	-0.097179
2013-01-04	-0.354580	0.568399	-0.887212	-1.719847
2013-01-05	1.948537	1.992504	0.514932	-0.492156

```
In [13]: df.index

Out[13]: DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
                        '2013-01-05', '2013-01-06'],
                        dtype='datetime64[ns]', freq='D')

In [14]: df.columns
```

```
Out[14]: Index(['A', 'B', 'C', 'D'], dtype='object')
```

```
In [15]: df.values
```

```
Out[15]: array([[ 0.88005897,  0.47766008,  0.65196433,  0.52581714],
                [-0.07069612,  0.81144008,  1.04258335,  1.42912146],
                [-0.41765574,  1.17158076,  0.88648631, -0.09717914],
                [-0.3545797 ,  0.56839862, -0.88721205, -1.71984699],
                [ 1.94853717,  1.99250448,  0.51493211, -0.49215609],
                [-0.67839866, -0.29133917, -3.33122589, -0.84379543]])
```

```
In [16]: # quick data summary
df.describe()
```

```
Out[16]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	0.217878	0.788374	-0.187079	-0.199673
std	1.005173	0.762630	1.687440	1.095553
min	-0.678399	-0.291339	-3.331226	-1.719847
25%	-0.401887	0.500345	-0.536676	-0.755886
50%	-0.212638	0.689919	0.583448	-0.294668
75%	0.642370	1.081546	0.827856	0.370068
max	1.948537	1.992504	1.042583	1.429121

```
In [17]: df.T
```

```
Out[17]:
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04	2013-01-05	2013-01-06
A	0.880059	-0.070696	-0.417656	-0.354580	1.948537	-0.678399
B	0.477660	0.811440	1.171581	0.568399	1.992504	-0.291339
C	0.651964	1.042583	0.886486	-0.887212	0.514932	-3.331226
D	0.525817	1.429121	-0.097179	-1.719847	-0.492156	-0.843795

```
In [18]: # axis 0 is index, axis 1 is columns
df.sort_index(axis=1, ascending=False)
```

```
Out[18]:
```

	D	C	B	A
2013-01-01	0.525817	0.651964	0.477660	0.880059
2013-01-02	1.429121	1.042583	0.811440	-0.070696
2013-01-03	-0.097179	0.886486	1.171581	-0.417656
2013-01-04	-1.719847	-0.887212	0.568399	-0.354580
2013-01-05	-0.492156	0.514932	1.992504	1.948537
2013-01-06	-0.843795	-3.331226	-0.291339	-0.678399

```
In [19]: # can sort by values too
df.sort_values('B')
```

```
Out[19]:
```

	A	B	C	D
2013-01-06	-0.678399	-0.291339	-3.331226	-0.843795
2013-01-01	0.880059	0.477660	0.651964	0.525817
2013-01-04	-0.354580	0.568399	-0.887212	-1.719847
2013-01-02	-0.070696	0.811440	1.042583	1.429121
2013-01-03	-0.417656	1.171581	0.886486	-0.097179
2013-01-05	1.948537	1.992504	0.514932	-0.492156

5.0.1 Selection

```
In [20]: # select a column (yields a series)
df['A']
```

```
Out[20]: 2013-01-01    0.880059
         2013-01-02   -0.070696
         2013-01-03   -0.417656
         2013-01-04   -0.354580
         2013-01-05    1.948537
         2013-01-06   -0.678399
         Freq: D, Name: A, dtype: float64
```

```
In [21]: # column names also attached to the object
df.A
```

```
Out[21]: 2013-01-01    0.880059
         2013-01-02   -0.070696
         2013-01-03   -0.417656
         2013-01-04   -0.354580
         2013-01-05    1.948537
         2013-01-06   -0.678399
         Freq: D, Name: A, dtype: float64
```

```
In [22]: # slicing works
df[0:3]
```

```
Out[22]:
```

	A	B	C	D
2013-01-01	0.880059	0.477660	0.651964	0.525817
2013-01-02	-0.070696	0.811440	1.042583	1.429121
2013-01-03	-0.417656	1.171581	0.886486	-0.097179

```
In [23]: df['20130102':'20130104']
```

```
Out[23]:
```

	A	B	C	D
2013-01-02	-0.070696	0.811440	1.042583	1.429121
2013-01-03	-0.417656	1.171581	0.886486	-0.097179
2013-01-04	-0.354580	0.568399	-0.887212	-1.719847

```
In [24]: # cross-section using a label
df.loc[dates[0]]
```

```
Out[24]: A    0.880059
         B    0.477660
         C    0.651964
         D    0.525817
         Name: 2013-01-01 00:00:00, dtype: float64
```

```
In [25]: # getting a scalar value
df.loc[dates[0], 'A']
```

Out [25]: 0.8800589662268992

```
In [26]: # select via position
df.iloc[3]
```

Out [26]:

A	-0.354580
B	0.568399
C	-0.887212
D	-1.719847

Name: 2013-01-04 00:00:00, dtype: float64

```
In [27]: df.iloc[3:5,0:2]
```

Out [27]:

	A	B
2013-01-04	-0.354580	0.568399
2013-01-05	1.948537	1.992504

```
In [28]: # column slicing
df.iloc[:,1:3]
```

Out [28]:

	B	C
2013-01-01	0.477660	0.651964
2013-01-02	0.811440	1.042583
2013-01-03	1.171581	0.886486
2013-01-04	0.568399	-0.887212
2013-01-05	1.992504	0.514932
2013-01-06	-0.291339	-3.331226

```
In [29]: # get a value by index
df.iloc[1,1]
```

Out [29]: 0.8114400831339718

```
In [30]: # boolean indexing
df[df.A > 0]
```

Out [30]:

	A	B	C	D
2013-01-01	0.880059	0.477660	0.651964	0.525817
2013-01-05	1.948537	1.992504	0.514932	-0.492156

```
In [31]: df[df > 0]
```

Out [31]:

	A	B	C	D
2013-01-01	0.880059	0.477660	0.651964	0.525817
2013-01-02	NaN	0.811440	1.042583	1.429121
2013-01-03	NaN	1.171581	0.886486	NaN
2013-01-04	NaN	0.568399	NaN	NaN
2013-01-05	1.948537	1.992504	0.514932	NaN
2013-01-06	NaN	NaN	NaN	NaN

```
In [32]: # filtering
df3 = df.copy()
df3['E'] = ['one', 'one', 'two', 'three', 'four', 'three']
print(df3['E'].isin(['two', 'four']))
df3[df3['E'].isin(['two', 'four'])]
```

```
2013-01-01    False
2013-01-02    False
2013-01-03     True
2013-01-04    False
2013-01-05     True
2013-01-06    False
Freq: D, Name: E, dtype: bool
```

```
Out[32]:
```

	A	B	C	D	E
2013-01-03	-0.417656	1.171581	0.886486	-0.097179	two
2013-01-05	1.948537	1.992504	0.514932	-0.492156	four

```
In [33]: # setting examples
df.at[dates[0], 'A'] = 0
df.iat[0,1] = 0
df.loc[:, 'D'] = np.array([5] * len(df))
df
```

```
Out[33]:
```

	A	B	C	D
2013-01-01	0.000000	0.000000	0.651964	5
2013-01-02	-0.070696	0.811440	1.042583	5
2013-01-03	-0.417656	1.171581	0.886486	5
2013-01-04	-0.354580	0.568399	-0.887212	5
2013-01-05	1.948537	1.992504	0.514932	5
2013-01-06	-0.678399	-0.291339	-3.331226	5

```
In [34]: # dealing with missing data
df4 = df.reindex(index=dates[0:4], columns=list(df.columns) + ['E'])
df4.loc[dates[0]:dates[1], 'E'] = 1
df4
```

```
Out[34]:
```

	A	B	C	D	E
2013-01-01	0.000000	0.000000	0.651964	5	1.0
2013-01-02	-0.070696	0.811440	1.042583	5	1.0
2013-01-03	-0.417656	1.171581	0.886486	5	NaN
2013-01-04	-0.354580	0.568399	-0.887212	5	NaN

```
In [35]: # drop rows with missing data
df4.dropna(how='any')
```

```
Out[35]:
```

	A	B	C	D	E
2013-01-01	0.000000	0.000000	0.651964	5	1.0
2013-01-02	-0.070696	0.811440	1.042583	5	1.0

```
In [36]: # fill missing data
df4.fillna(value=5)
```

```
Out[36]:
```

	A	B	C	D	E
2013-01-01	0.000000	0.000000	0.651964	5	1.0
2013-01-02	-0.070696	0.811440	1.042583	5	1.0
2013-01-03	-0.417656	1.171581	0.886486	5	5.0
2013-01-04	-0.354580	0.568399	-0.887212	5	5.0

```
In [37]: # boolean mask for nan values
pd.isnull(df4)
```

```
Out[37]:
```

	A	B	C	D	E
2013-01-01	False	False	False	False	False
2013-01-02	False	False	False	False	False
2013-01-03	False	False	False	False	True
2013-01-04	False	False	False	False	True

5.0.2 Operations

```
In [38]: df.mean()
```

```
Out[38]: A    0.071201
B    0.708764
C   -0.187079
D    5.000000
dtype: float64
```

```
In [39]: # pivot the mean calculation
df.mean(1)
```

```
Out[39]: 2013-01-01    1.412991
2013-01-02    1.695832
2013-01-03    1.660103
2013-01-04    1.081652
2013-01-05    2.363993
2013-01-06    0.174759
Freq: D, dtype: float64
```

```
In [40]: # aligning objects with different dimensions
s = pd.Series([1,3,5,np.nan,6,8],index=dates).shift(2)
df.sub(s,axis='index')
```

```
Out[40]:
```

	A	B	C	D
2013-01-01	NaN	NaN	NaN	NaN
2013-01-02	NaN	NaN	NaN	NaN
2013-01-03	-1.417656	0.171581	-0.113514	4.0
2013-01-04	-3.354580	-2.431601	-3.887212	2.0
2013-01-05	-3.051463	-3.007496	-4.485068	0.0
2013-01-06	NaN	NaN	NaN	NaN


```
In [41]: # applying functions
df.apply(np.cumsum)
```

```
Out[41]:
```

	A	B	C	D
2013-01-01	0.000000	0.000000	0.651964	5
2013-01-02	-0.070696	0.811440	1.694548	10
2013-01-03	-0.488352	1.983021	2.581034	15
2013-01-04	-0.842932	2.551419	1.693822	20
2013-01-05	1.105606	4.543924	2.208754	25
2013-01-06	0.427207	4.252585	-1.122472	30

```
In [42]: df.apply(lambda x: x.max() - x.min())
```

```
Out[42]: A    2.626936
         B    2.283844
         C    4.373809
         D    0.000000
         dtype: float64
```

```
In [43]: # simple count aggregation
s = pd.Series(np.random.randint(0,7,size=10))
s.value_counts()
```

```
Out[43]: 2    4
         4    2
         0    2
         3    1
         1    1
         dtype: int64
```

5.0.3 Merging / Grouping / Shaping

```
In [44]: # concatenation
df = pd.DataFrame(np.random.randn(10, 4))
pieces = [df[:3], df[3:7], df[7:]]
pd.concat(pieces)
```

```
Out[44]:
```

	0	1	2	3
0	-0.822660	-0.371737	0.017555	-0.324006
1	-1.247939	-0.247542	0.350824	-0.414374
2	1.005921	-0.426244	-1.806617	0.392695
3	-0.589945	0.503436	1.018880	0.296445
4	1.110050	-0.989417	1.553230	-0.988084
5	0.393925	-0.860651	-0.348174	0.271126
6	0.711039	-0.308275	-0.048588	1.395454
7	0.055404	-0.676232	1.261214	0.140917
8	-1.423938	-0.428276	-0.209122	-1.004157
9	0.564595	1.225142	-0.732768	-0.786322

```
In [45]: # SQL-style join
left = pd.DataFrame({'key': ['foo', 'foo'], 'lval': [1, 2]})
right = pd.DataFrame({'key': ['foo', 'foo'], 'rval': [4, 5]})
pd.merge(left, right, on='key')
```

```
Out[45]:
```

	key	lval	rval
0	foo	1	4
1	foo	1	5
2	foo	2	4
3	foo	2	5

```
In [46]: # append
df = pd.DataFrame(np.random.randn(8, 4), columns=['A', 'B', 'C', 'D'])
s = df.iloc[3]
df.append(s, ignore_index=True)
```

```
Out[46]:
```

	A	B	C	D
0	-0.348170	-0.756201	0.833933	-0.513968
1	0.966034	-0.844801	-1.263234	-0.632908
2	-1.694731	-0.345543	2.104304	0.182621
3	-0.214822	1.075815	0.167778	0.836616
4	-0.413597	-0.256741	-0.803030	0.265014
5	-0.607043	-0.058836	0.554306	-0.029471
6	-1.828674	1.318696	0.541203	0.264778
7	-0.510968	1.198982	-0.139671	-0.997434
8	-0.214822	1.075815	0.167778	0.836616

```
In [47]: df = pd.DataFrame(
    { 'A' : ['foo', 'bar', 'foo', 'bar', 'foo', 'bar', 'foo', 'foo'],
      'B' : ['one', 'one', 'two', 'three', 'two', 'two', 'one', 'three'],
      'C' : np.random.randn(8),
      'D' : np.random.randn(8) })
df
```

```
Out[47]:
```

	A	B	C	D
0	foo	one	-1.723702	1.106293
1	bar	one	1.506323	1.236936
2	foo	two	0.847108	0.814905
3	bar	three	-1.316381	-1.212962
4	foo	two	0.341018	0.174246
5	bar	two	-0.025023	1.441938
6	foo	one	0.367602	-2.741238
7	foo	three	-0.038510	-0.871220

```
In [48]: # group by
df.groupby('A').sum()
```

```
Out[48]:
```

	C	D
A		
bar	0.164919	1.465912
foo	-0.206483	-1.517014

```
In [49]: # group by multiple columns
df.groupby(['A', 'B']).sum()
```

```
Out[49]:
```

		C	D
A	B		
bar	one	1.506323	1.236936
	three	-1.316381	-1.212962
	two	-0.025023	1.441938
foo	one	-1.356100	-1.634945
	three	-0.038510	-0.871220
	two	1.188127	0.989151

```
In [50]: df = pd.DataFrame(
    { 'A' : ['one', 'one', 'two', 'three'] * 3,
      'B' : ['A', 'B', 'C'] * 4,
      'C' : ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 2,
      'D' : np.random.randn(12),
      'E' : np.random.randn(12)} )
df
```

```
Out[50]:
```

		A	B	C	D	E
0	one	A	foo	0.077271	1.094401	
1	one	B	foo	0.396550	-0.512023	
2	two	C	foo	-1.480050	0.045612	
3	three	A	bar	1.927040	-1.284882	
4	one	B	bar	-0.397431	-0.566343	
5	one	C	bar	-1.188633	1.838150	
6	two	A	foo	-0.684879	0.033144	
7	three	B	foo	2.479158	1.126320	
8	one	C	foo	-0.023391	0.934111	
9	one	A	bar	-0.194993	0.576714	
10	two	B	bar	-0.579569	2.624459	
11	three	C	bar	0.047084	0.941124	

```
In [51]: # pivot table
pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
```

```
Out[51]:
```

		C	bar	foo
A	B			
one	A	-0.194993	0.077271	
	B	-0.397431	0.396550	
	C	-1.188633	-0.023391	
three	A	1.927040	NaN	
	B	NaN	2.479158	
	C	0.047084	NaN	
two	A	NaN	-0.684879	
	B	-0.579569	NaN	
	C	NaN	-1.480050	

5.0.4 Time Series

```
In [52]: # time period resampling
rng = pd.date_range('1/1/2012', periods=100, freq='S')
ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)
ts.resample('5Min', how='sum')
```

/home/nareshguru77/anaconda2/envs/py36/lib/python3.6/site-packages/ipykernel_launcher.py:4: FutureWarning: the new syntax is .resample(...).sum()
after removing the cwd from sys.path.

```
Out[52]: 2012-01-01    24754
         Freq: 5T, dtype: int64
```

```
In [53]: rng = pd.date_range('1/1/2012', periods=5, freq='M')
ts = pd.Series(np.random.randn(len(rng)), index=rng)
ts
```

```
Out[53]: 2012-01-31    -0.863508
         2012-02-29    -1.921387
         2012-03-31     0.329550
         2012-04-30     1.271644
         2012-05-31     0.045044
         Freq: M, dtype: float64
```


```
In [54]: ps = ts.to_period()
         ps.to_timestamp()
```

```
Out[54]: 2012-01-01    -0.863508
         2012-02-01    -1.921387
         2012-03-01     0.329550
         2012-04-01     1.271644
         2012-05-01     0.045044
         Freq: MS, dtype: float64
```

5.0.5 Plotting

```
In [55]: # time series plot
ts = pd.Series(np.random.randn(1000), index=pd.date_range('1/1/2000', periods=1000))
ts = ts.cumsum()
ts.plot()
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5fe776c2e8>
```




In [56]: *# plot with a data frame*

```
df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index, columns=['A', 'B', 'C', 'D'])  
df = df.cumsum()  
plt.figure(); df.plot(); plt.legend(loc='best')
```

Out[56]: <matplotlib.legend.Legend at 0x7f5fe3bfc38>

<matplotlib.figure.Figure at 0x7f5fe3b932b0>



5.0.6 Input/Output

```
In [57]: # write to a csv file
df.to_csv('foo.csv', index=False)
```

```
In [58]: # read file back in
path = './foo.csv'
newDf = pd.read_csv(path)
newDf.head()
```

```
Out[58]:
```

	A	B	C	D
0	2.005550	-0.753783	-1.231043	-0.508113
1	1.998454	-0.858450	-0.598276	-0.647307
2	1.260837	-0.882566	0.206137	-2.050626
3	1.496103	-0.480195	0.055814	-2.499547
4	1.474878	-0.546122	-1.136639	-2.883347

```
In [59]: # remove the file
import os
os.remove(path)
```

```
In [60]: # can also do Excel
df.to_excel('foo.xlsx', sheet_name='Sheet1')
```

```
In [61]: newDf2 = pd.read_excel('foo.xlsx', 'Sheet1', index_col=None, na_values=['NA'])
newDf2.head()
```

```
Out [61]:
```

	A	B	C	D
2000-01-01	2.005550	-0.753783	-1.231043	-0.508113
2000-01-02	1.998454	-0.858450	-0.598276	-0.647307
2000-01-03	1.260837	-0.882566	0.206137	-2.050626
2000-01-04	1.496103	-0.480195	0.055814	-2.499547
2000-01-05	1.474878	-0.546122	-1.136639	-2.883347

```
In [62]: os.remove('foo.xlsx')
```